

Linux Kernel Programming

Lab 04: Debugging in the Linux Kernel

AUTHOR

Redha Gouicem & Julien Sopena

In this lab, you will familiarize yourself with the kernel debugging tool *kgdb* through examples of buggy modules. You will also come back to the proper use of lists in the kernel.

Task 1: Connecting the *kgdb* debugger

The Linux kernel embeds a debugger that can be accessed through various interfaces. One of them, *kgdb*, allows users to debug the kernel from a different machine (e.g., the host machine in our virtualized setup) using *gdb*.

From now on, you will run *gdb* on the host machine, connect it to the Linux kernel running in the guest VM, and use it to debug kernel code.

Question 1

Using one of the configuration tools of the kernel (*make menuconfig*, etc), check that *kgdb* and *kdb* are enabled, and enable them if they are not.

Tip

Most of the configuration options we will use in this exercise are located in the *Kernel hacking* section.

You should also use the search function to find where some options are located and their dependencies.

Question 2

When the kernel crashes, *kgdb* is triggered in the VM, which enables users to connect to it with *gdb*. To do so, you need to configure which serial port *kgdb* should use to connect with *gdb*. Furthermore, this serial port needs to be accessible from the host machine, which is made possible by QEMU by exporting a serial port to a file, a socket, etc.

Open the *qemu-run-externKernel.sh* script and figure out how the serial port for *kgdb* is exported.

Question 3

The kernel command line option *kgdbwait* makes the kernel pause during the boot process to wait for a debugger to connect to it. Use this option (check out the *qemu-run-externKernel.sh* script) and start a VM. Start *gdb* on your host machine on the *vmlinux* binary available at the root of your kernel directory. Finally, when the VM is paused, connect *gdb* to your guest kernel with the command *target remote* and the serial port exported.

Question 4

In *kgdb*, what is the difference between the commands `info thread` and `monitor ps`? Use the commands `help` and `monitor help` in *gdb* to figure it out.

Question 5

What does the `continue` command do?

Question 6

You can give control back to *kgdb* without the kernel crashing or pausing on purpose with *Magic SysRq Key*. First, make sure your kernel is compiled with the option `CONFIG_MAGIC_SYSRQ` enabled.

To interrupt the VM and allow *kgdb* to take control, you can use the combination of keys Alt+SysRq+g, or simulate this by writing 'g' in the `/proc/sysrq-trigger` file.

The *Magic SysRq Key* mechanism offer powerful features that may simplify debugging. You can find all these features on the [Magic SysRq key Wikipedia page](#).

Task 2: Using *kgdb*

Question 1

The global structure `init_uts_ns` stores the information returned by `uname -r`. In the kernel source, find where this structure is defined and initialized and check its members.

Question 2

Using the `print` and `set variable` commands in *gdb*, display the content of the structure and modify it so that `uname -r` returns another name.

Task 3: My first bug

Now, we will provide two modules, bugged on purpose. The goal of this task is not to find the bugs by analyzing the code but by using *kgdb* instead.

Question 1

Check out the sources of the module `hanging` and explain its behavior.

Question 2

Compile the `hanging` module for your guest kernel, copy it into your VM (with the shared directory, ssh or in the *home* disk image), and load it in the VM. Wait for a while (~30 secondes). What happens? Does your kernel crash?

Analyze the output of the message printed in the console.

Question 3

In order to trigger *kgdb* when the message appears, modify the configuration of your kernel to trigger a *panic* in that case. Recompile your kernel and module, and load it.

Question 4

In *gdb*, dump the backtrace (with the `backtrace` command, or `bt`) when the *panic* happens. Does it match the code of your module?

Question 5

Use the commands `monitor ps` and `monitor btp` to dump the backtrace of the module correctly.

Question 6

In *kgdb*, check that your module is properly loaded with `monitor lsmod`. The first address displayed is the address of the `struct module` of the module. What is the second address?

Question 7

How can you solve the bug? Propose two solutions:

- one that modifies only the code of the module without changing the duration of the hanging (the schedule timeout);
- one that only modifies the configuration of the kernel.

Task 4: Dynamic debug printing

The kernel provides *dynamic debugging* that allows to dynamically enable/disable the debug prints at run time (`printk()` with `KERN_DEBUG` or `pr_debug()`). This is a great tool if you want a large number of debug prints without always impacting the performance of your code (yes, printing is costly).

Documentation: `Documentation/admin-guide/dynamic-debug-howto.rst`

Question 1

The module `prdebug` provided periodically prints the number of interrupts received by core 0.

Check out the sources of the module and try to understand its behavior. Compile it and test it. Why aren't there any prints in the kernel logs?

Question 2

The file `/sys/kernel/debug/dynamic_debug/control` allows you to configure the behavior of dynamic printing.

Load the `prdebug` module into the VM and enable the prints generated by the module with the following command:

```
echo -n 'module prdebug +p' > /sys/kernel/debug/dynamic_debug/control
```

Explain this command with the help of the documentation.

Question 3

You can dynamically add information to your debug prints, such as the name of the module, the function or the line number of the `pr_debug()`.

Add these information to your debug prints by sending a command to the `control` file used in the previous question.

Question 4

You can also apply filters and only show some `pr_debug()` of your module, e.g., *only the print on line 13*.

Unload and reload your module in the VM. Check that the dynamic debugging configuration has been reset.

With the help of the documentation, enable **only** the message that shows the number of received interrupts.

Task 5: Debugging a module

In this task, multiple bugs have been introduced in the provided module `kcpustat`. The goal of this task is to detect, analyze and fix these bugs with debugging options of the kernel and `kgdb`.

Warning

Do not modify the sources of the provided module if not explicitly asked to.

Question 1

The module `kcpustat` collects CPU usage statistics (like the command `time`) and displays them in the kernel logs.

Check out the sources of the module and try to understand how it works.

Question 2

Compile the module and load it in the VM. Wait for a few seconds. What happens?

Question 3

Connect `gdb` to your VM and check the results of `backtrace` and `monitor btp`. The results is hard to exploit...

With the commands `monitor lsmod` and `add-symbol-file`, load the symbols of the module (the `.o` file) at its load address.

Question 4

Which line of the module's code triggers the fault? Fix the bug without deleting the faulty line of code.

Question 5

Try your fixed module. What happens? Again, connect `gdb` to figure out which data structure causes the problem.

Question 6

The kernel has options to help debug some data structures, which can be enabled in the kernel configuration. Enable this option for the data structure you want to debug.

You should now have new information in the kernel console when the bug happens. What do you deduce from this information?

Propose a fix for the bug.

Question 7

The module also doesn't properly manage its memory. The kernel provides a memory leak detector that you can enable in the kernel configuration.

Enable it and read its documentation.

Question 8

Use the kernel memory leak detector to find where these leaks happen and fix them.